Trustworthy AI Agents:

Toward Safe, Secure and Sovereign Agents

August 18, 2026
Dr. Mark Maybury, Editor

Call for Book Chapter Submissions

As AI models increase in capability, agency and autonomy, their impact on individuals and society grows. This interdisciplinary book shares state of the art results from researchers, practitioners, policy makers and students from diverse fields including computer science, sociology, philosophy, and law who collaborated to create a roadmap considering the scientific and technical, security, economic, environmental, regulatory, social and ethical implications of Agentic AI to advance responsible if not beneficent agents and their enabling ecosystem. Following the IJCAI-ECAI 2026 Workshop on Safe Agents and public Panel, this new scientific collection will include a mix of invited and submitted papers designed to create an interdisciplinary research foundation that can help advance agent progress responsibly. This global topic is timely, interdisciplinary, and necessary, addressing a critical need in the AI community and society at large.

Submissions:
We invite submissions that include surveys of existing agent gaps or forecasts, late breaking insights into future safety and security of agents and speculative but clearly articulated responsible agent futures. Submissions will be evaluated based on criteria of novelty, originality, technical correctness, clarity, significance of results, and potential impact. Given the roadmap focus of the book, we will encourage results that both carefully articulated gaps, evidence-based forecasts as well as visionary but clearly articulated desired future capabilities or outcomes. The final formal of the book will be in single column so to minimize formatting changes and speed publication, so proposed chapter submissions should be made to this NEW book chapter review web site:

https://chairingtool.com/conferences/safeagents/main-track?role=author

in Springer's format using the MS Word templates and the information for authors available including a required license to publish form, all of which can be found here:

https://link.springer.com/series/558/information-for-authors-and-editors

The corresponding author, who must match the corresponding author marked on the paper, must have the full right, power, and authority to sign the agreement on behalf of all of the authors of a particular paper, and accepts responsibility for releasing this material on their behalf. Authors much adhered to this Code of Conduct (https://www.springernature.com/gp/authors/book-authors-code-of-conduct).

For any additional material (e.g., open source code, data or supporting materials), we will utilize the GitHub book web site: **github.com/mtmaybury/SafeAgentsBook**

The paper length is SINGLE column 12-15 (or more) pages for long chapters and 6-11 pages for short chapters to allow authors to add in new results and/or insights from the workshop and should include a short summary at the end of the prosed chapter and a call out box that points to practical tools or on line sources as well as a set of discussion questions. We will extend submissions to others who have papers of high relevance to the topic. Prior acceptance at the workshop does NOT guarantee inclusion in the book as submissions will go through a fresh round of peer review. If you know of relevant new research that you believe would add value to this collection on Safe, Secure and Sovereign Agents, please share their contact information and share this CFP with them. Given the expert focus of this research, each chapter contribution will be responsible for review of 3 others papers. If authors submit an additional paper of distinct research, they will be required to review an additional 3 papers for each new submission. An editorial review board of the workshop organizers will provide overall quality assurance of the book beyond peer review. A draft overview of the workshop roadmap is now available now on the GitHub site and will be extended based on reviewer feedback and extended and new chapter submissions.

If any author contributing to Safe, Secure and Sovereign Agents is interested in Open Choice (which applies to individual papers), please refer to Springer's webpage for prices and additional information. Springer requires the invoicing address and the CC-BY licence-to-publish agreement at least three weeks before the deadline for the files.

Submissions do not need to be anonymized. To increase the value to readers, submissions should include two additional items to each submitted chapter (no more than one page each, can be as short as a paragraph):

1. **Boxed "From Theory to Practice"** sidebar with open-source links and/or code snippets to guide the reader to how they can benefit from your work. You can store this information in our workshop GitHub if that is more convenient.
2. **End-of-chapter summary** and **discussion questions** so that students, researchers and faculty can use the book in their practice.

Deadlines:

Book Chapter Submission: October 5, 2026
Reviewing Period (Each submitting chapter must peer review 3 other submissions): Oct 5 to Nov 5, 2026

SPC Review Nov 6-16 2026.
Discussion Phase: Nov 17-Nov 23

Book Chapter Notification: November 24, 2026
Camera Ready Copy: December 14, 2027
Draft Book: 1 January 2028.

Editorial Review Board

- Dr. Mark Maybury, VP Commercialization, Lockheed Martin; Board Director, Advanced Cyber Security Center. mark.t.maybury@lmco.com (chair and primary contact)
- Prof. Wolfgang Wahlster, Founding Director and CEA, DFKI. Wolfgang.Wahlster@dfki.de
- Ms. Josephine Liu, Chair, APAC Tech Policy Committee, Association for Computing Machinery (ACM), yuyinl@acm.org
- Dr. Mihai Christodorescu, Research Scientist, Security and Privacy, Google, christodorescu@google.com
- Dr. Stuart Battersby, sbattersby@redhat.com, Open Source AI Security Lead, Red Hat
- Dr. Francesca Rossi, IBM Fellow, IBM Global Leader for Responsible AI and AI Governance, T.J. Watson Research Center, Yorktown Heights, USA Francesca.Rossi2@ibm.com

Goal and Product: A book of state of the art papers the contribute that articulate or amplify elements of a Trustworthy Agentic AI Roadmap (TAP) articulating, current *critical gaps* or *challenges*, planned key *capabilities* and desired *outcomes*. The book will include an introductory chapter including an update of the Safe Agent roadmap created at the 2026 workshop.

Critical Questions of Interest: These include but are not limited to:

- **Agentic AI**: Agentic AI, or goal-seeking AI systems that can act autonomously to achieve specific objectives are expanding in capability and application. While agents promise to revolutionize everything from education to customer service to healthcare, they introduce new challenges such as control (e.g., delegation of authority, mixed-initiative), accountability, and verification and validation of evolving learning systems as well as perils arising from bias, brittleness and belligerence [Maybury 2025a,b] and resultant harms (e.g., financial, physical, social, reputational). How do we ensure and assure progress and avoid unintended and undesirable consequences?
- **Models and Autonomy**: The global deployment and use of large language/process models, small models for edge devices, and the introduction of large world models is accelerating new cases and advancing agents and autonomous systems across the enterprise and at the edge. Their computational and energy appetite for model training and inference drive a need for resource-efficient infrastructure and processing. How can we accelerate the state of the art in algorithms, architectures, interoperability (e.g., Model Context Protocol) and security to simultaneously advance performance, usability, prosperity, transparency, and safety?
- **Society, Safety and Sustainability.** Current LLMs suffer from being biased, brittle, and baroque with risk of significant harms. A rapidly expanding number of agents and humanoid robots add new risks of belligerence. Beginning to address these concerns, societies have introduced national guidelines (e.g., NIST AI RMF), regulations (e.g., DoD 3000.09, Chinese Regulation on the Governance of Generative AI), strategies (e.g., Germany's AI Action Plan, US AI Action Plan, UK AI Strategy), laws (e.g., EU AI Act, GDPR), and sector-specific safety and security guidelines. What are effective methods for responsible development, deployment and operation that can mitigate risks to society? How can governance frameworks, standards, and regulatory mechanisms be effectively implemented and coordinated across jurisdictions and institutions to ensure responsible development and deployment of agentic AI systems?

Book Themes

Based on final accepted submissions, the book will be structured around several key sub-themes of agent safety to ensure a comprehensive and focused discussion including:

- **Ethical Foundations and Governance**: theoretical and philosophical underpinnings of ethical AI for agents. Topics can include:
  - Assessing current and new ethical frameworks for autonomous and agentic systems.
  - The role of law, regulation and policy in guiding agent development.
  - Philosophical debate around agent consciousness, rights and responsibilities.

- **Accuracy, Bias, Fairness, Transparency**: A core technical and social component of the workshop will address:
  - Methods for improving agent accuracy and self-awareness (e.g., limitations, confidence).
  - Successful approaches to detecting and mitigating dangers such as bias and hallucination in large datasets and models.
  - Techniques for creating more transparent and interpretable models and agents.

- Evaluating the impact of agents on different demographic groups.

- **Human-Agent Collaboration and Trust**: As intelligent agents becomes more integrated into daily life, understanding the dynamics of human-agent collaboration is crucial, including:
  - Designing intuitive and trustworthy user interfaces for agentic AI.
  - The role of agents in decision support systems and associated risks (e.g., awareness, over/under confidence, deskilling).
  - Governance frameworks for autonomous decision-making by AI agents, including responsibility allocation, oversight mechanisms, and human-in-the-loop or human-on-the-loop models.
  - Strategies for building public trust and understanding of agent technologies.

- **Real-world Applications and Case Studies**: Case studies from various industries such as:
  - Responsible development and deployment in low to high risk use cases
  - Responsible agents in high assurance applications with human, financial, and security risks.
  - Ethical and operational challenges in using agents in high assurance and regulated industries such as healthcare, disaster management, pharmaceutical, financial services, telecommunications, logistics, energy, agriculture, defense, and space.
  - The use of AI in creative fields and its impact on human learning, training, and creativity.

Target Book Chapter Contributors

- **Academic Researchers**: Computer scientists, ethicists, sociologists, and legal scholars.
- **Industry Practitioners**: AI engineers, product managers, and data scientists working on real-world applications.
- **Policymakers/Regulators**: Officials and orgs concerned with AI/Agent governance and safety.
- **Students**: Graduate and undergraduate students interested in social and ethical dimensions of AI.

Dr. Mark Maybury
Chair, Safe Agents Workshop, IJCAI
mark.t.maybury@lmco.com
mtmaybury@gmail.com